

**M340L-CS, SPRING 2024**  
**Matrices and Matrix Computations**

**Project: QR decomposition**

**Theory.** We have seen an application of the QR decomposition to least square problems. Later, we will also use it as a part of many other numerical linear algebra algorithms. It is therefore useful to have some *efficient* and *numerically stable* algorithms.

The focus of this project is on the stability part. We will compare three different QR decomposition algorithms: Gram–Schmidt, modified Gram–Schmidt, Householder reflections. If  $A$  is an  $n \times k$  matrix, and the program outputs an  $n \times k$  matrix  $Q$  and a  $k \times k$  matrix  $R$ , then we want

$$Q^T Q = I_k, \quad A = QR,$$

so two good measures of error are

$$\varepsilon_{\perp} = \|Q^T Q - I_k\|_{\max}, \quad \varepsilon_s = \|A - QR\|_{\max}$$

where  $\|X\|_{\max}$  for a matrix  $X$  just means the largest absolute value among all of its entries.

**Problems.**

1. Write a program to perform Gram–Schmidt on an  $n \times k$  matrix  $A$ , i.e. perform the following procedure for  $i = 1, \dots, k$ .

$$\begin{aligned} \mathbf{v}_i &\leftarrow \mathbf{v}_i - (\mathbf{q}_1 \cdot \mathbf{v}_i)\mathbf{q}_1 - \dots - (\mathbf{q}_{i-1} \cdot \mathbf{v}_i)\mathbf{q}_{i-1} \\ \mathbf{q}_i &\leftarrow \mathbf{v}_i / \|\mathbf{v}_i\| \end{aligned}$$

The matrix  $Q$  is  $[\mathbf{q}_1 \ \dots \ \mathbf{q}_k]$ . Compute  $R$  using  $R = Q^T A$ . Output the matrices  $Q$  and  $R$ , as well as the two errors  $\varepsilon_{\perp}$ ,  $\varepsilon_s$ .

(Make sure you are not accidentally doing the modification described in the following part. This could happen from passing by reference.)

2. In a different function, make the following modification to the iteration step

$$\begin{aligned} &\text{for } j = 1, \dots, i-1 : \\ &\quad \mathbf{v}_i \leftarrow \mathbf{v}_i - (\mathbf{q}_j \cdot \mathbf{v}_i)\mathbf{q}_j \\ &\quad \mathbf{q}_i \leftarrow \mathbf{v}_i / \|\mathbf{v}_i\| \end{aligned}$$

In other words, instead of computing all of the coefficients at once, we compute them one at a time, updating  $\mathbf{v}_i$  before each computation. The output should consist of the same quadruple  $(Q, R, \varepsilon_{\perp}, \varepsilon_s)$  as in part 1.

3. Recall the Householder reflection matrix

$$R = I_n - \frac{2\mathbf{w}\mathbf{w}^T}{\mathbf{w}^T \mathbf{w}}$$

This is not a sparse matrix in general. Despite this, describe a way to compute  $RA$  for an arbitrary  $A$  with  $O(n^2)$  operations (the usual way would take  $O(n^3)$  operations).

4. Using this, implement the Householder reflection method for QR factorization. See the supplementary notes for Lecture 21 for details. To compute  $Q$ , you might find it helpful to initialize  $B = I_n$ , and every time you apply a reflection to the input  $A$ , also do it to  $B$ . The result is  $Q^T$ . You are recommended to do the calculations in place, overwriting entries of the input (or a deep copy of the input).

Its output is the same quadruple as before, so if the input is not a square matrix, you should truncate  $Q$  and  $R$  as described in Homework 7, Q3.

- How many operations do the three methods take on an  $n \times k$  matrix? Give your answer using the big  $O$  notation, ignoring constants.

The following questions will test your three programs on both random matrices and the infamous Hilbert matrices, giving you a sense of their numerical stability.

- Run the three programs on some  $200 \times 200$  matrices whose entries are randomly selected between  $-1$  and  $1$ . Write down the average of  $\varepsilon_{\perp}$ . How does this average compare?
- You should have gotten reasonably accurate results, but this is not necessarily a good model for matrices showing up in real life. The Hilbert matrix is the  $n \times n$  matrix with entries

$$H_{ij} = \frac{1}{i+j-1}, \quad i, j = 1, \dots, n$$

This is a famous example of an *ill-conditioned* matrix, meaning it tends to drastically amplify round off errors, which makes numerical algorithms challenging. Since we are normalizing the columns of  $Q$ , the maximum possible value of  $\varepsilon_{\perp}$  is  $1$ . In this case, the supposedly orthogonal  $Q$  is close to being singular, and it is completely useless for any application.

For each of the three methods, find the smallest  $n$  such that  $\varepsilon_{\perp} > 0.9999$  when applied to the  $n \times n$  Hilbert matrix. If this is taking too much time, just say you can't find an  $n$ .

- We introduce a regularized version of the Hilbert matrix, which is slightly better behaved. Let  $\varepsilon = 0.0001$ . Define

$$H_{\varepsilon} = H + \varepsilon I_n$$

Apply the three methods to  $H_{\varepsilon}$  for  $n \leq 100$ . Make a graph which shows how the errors (both  $\varepsilon_{\perp}$  and  $\varepsilon_s$ ) depend on  $n$ . It is useful to put logarithmic scale on the  $y$ -axis.

- Give some plausible reasons for the difference in errors across the three methods.

**Remarks/Hints.** It might be of interest to look up the *condition number* of a matrix. Roughly speaking, this measures the amount of distortion that happens when  $A$  is applied. Large distortion means a small error in the input becomes a noticeable error in the output. By one definition, the condition number of a matrix is the quotient of its largest and its smallest singular value. In many practical settings, the matrices used have a few large singular values, and most of the other singular values are very close to  $0$ , representing noise. This is one reason ill-conditioned matrices show up in actual problems. The regularization of adding  $\varepsilon I_n$  keeps the smallest singular value greater than  $\varepsilon$  without significantly modifying the larger ones, so its condition number is smaller.

The Hilbert matrix famously has a large condition number. On the other extreme, orthogonal matrices have condition number  $1$  (in the  $L^2$ -norm). Intuitively, orthogonal matrices are just (a sequence of) rotations or reflections, and a small error in any direction stays small, just in a possibly different direction. This is why in ill-conditioned problems, orthogonal matrices are preferred. This further explains why the singular value decomposition is so useful.

In practice, when using Householder reflections,  $Q$  is almost never stored as a matrix. It is usually stored as a sequence of reflection matrices, and each one only takes  $O(n)$ -space to store the vector  $\mathbf{w}$ . To multiply by  $Q$ , one uses the type of methods you found in part 6.